

DocMind: A Self-Hosted Retrieval-Augmented Generation Platform with Measurable Retrieval Quality, Deterministic Layered Caching, and Operations-Grade Governance

Xiuquan Yu

Independent project report — source code available from the author

August 30, 2026

Abstract

Retrieval-augmented generation (RAG) systems are usually demonstrated rather than measured: prototypes are wired to a vector store and a prompt, but provide no answer to how much a design change helps, which queries are silently dropped, or what a system costs to run. This report describes DocMind, a self-hosted multi-assistant RAG platform built around three engineering commitments. First, retrieval quality is *measured*: a fixed 47-question benchmark (30 canonical, 17 adversarially colloquial) gates every change, showing that hybrid BM25–dense fusion via reciprocal rank fusion (RRF) lifts adversarial Recall@4 from 94.1% to 100.0%, while cross-encoder reranking improves MRR by 7.8% at a deliberate one-hit recall cost imposed by threshold filtering. Second, cost and latency are *engineered*, not hoped for: a four-layer cache hierarchy keyed by determinism (tokenizer, document embeddings, query embeddings, and rerank results) removes all upstream API round-trips for repeated questions — throughput on a warm replayed-question workload rises from 3.1 to 801 QPS with P95 latency dropping from 2,573 ms to 11 ms, while cold queries show no regression. Third, fine-tuning is applied *narrowly*: a rank-8 LoRA adapter tuned on a 1.5B model rewrites colloquial queries before retrieval, improving overall Recall@4 by 3.97 points at zero marginal API cost; a controlled experiment attributes the remaining headroom to corpus coverage rather than rewriting quality. The platform further integrates evidence-grounded refusal, RAGAS-style generation metrics, document-level access control, an MCP server surface, and full audit and alerting, and is deployed as a single docker-compose stack behind an nginx entry point.

Keywords: retrieval-augmented generation; hybrid retrieval; reciprocal rank fusion; cross-encoder reranking; LoRA; query rewriting; caching; evaluation; self-hosting

1 Introduction

Production LLM applications in the retrieval-augmented generation (RAG) style are now straightforward to prototype: embed a corpus, retrieve top- k , and prompt a chat model. The hard part begins after the demo. Teams cannot answer basic questions — which fraction of user queries fails silently, how a chunking change affects recall, what a month of traffic costs, or whether an answer is grounded in retrieved evidence. We call this gap the difference between a *RAG demo* and a *RAG system*, and we built DocMind to close it.

DocMind is a self-hosted, multi-assistant, multi-knowledge-base platform. Its retrieval pipeline is hybrid (BM25 plus dense vectors fused by reciprocal rank fusion, followed by a cross-encoder reranker); its agent is a hand-written ReAct loop with an evidence-grounded refusal path; and its operation is wrapped in an evaluation harness, cost accounting, alerting,




Figure 1: DocMind deployment and data flow. Enhancements (reranker, caches, web search) degrade gracefully; the answering path never depends on them succeeding.

and audit logging. The design philosophy is deliberately conservative: every enhancement component must degrade gracefully and must never block the answering path, and every claimed improvement must come with a reproducible number from a fixed benchmark.

This report makes four concrete contributions.

1. A benchmark-driven comparison of three retrieval routes on canonical and adversarially colloquial question sets, with an analysis of why rerank thresholds trade recall for refusal safety (Section 3).
2. A four-layer cache hierarchy organized by the determinism of each cached function, with an A/B evaluation on identical code showing a $258\times$ throughput improvement on warm replayed-question workloads and zero regression on cold queries (Section 4).
3. A narrow-task fine-tuning study: a LoRA-adapted 1.5B query rewriter with deterministic training-data construction, plus a controlled experiment that separates rewriting quality from corpus coverage (Section 5).
4. An operations-grade governance layer — audit, alerting, backup, access control, and an MCP tool surface — that makes the platform deployable inside organizations (Section 7).

2 System Overview

Figure 1 shows the deployment topology. A single nginx process serves the React single-page application and reverse-proxies REST and server-sent-event (SSE) traffic to a FastAPI host; the whole stack ships as one docker-compose project behind ports 80/443. The application host embeds a hand-written ReAct agent — no agent framework — with a tool registry exposing knowledge retrieval, web search with a four-engine fallback chain, and MCP tools. Authentication supports local accounts (PBKDF2 with per-user salt), forced first-login password rotation, and optional enterprise LDAP. Document-level access control is enforced inside the retrieval pipeline itself: unauthorized documents are filtered *before* truncation so that restricted content neither leaks nor crowds out authorized candidates.

Two principles shaped the implementation. First, the agent core is owned code: the ReAct loop, tool-result sanitization, injection guards, and the anti-dead-loop budget are implemented in-tree, which keeps the trust boundary explicit and the behavior auditable. Second, every enhancement is optional at runtime: the reranker falls back to RRF ordering on failure, Langfuse

Table 1: Retrieval routes on the 47-question benchmark (30 canonical, 17 adversarial). Every number is produced by the benchmark script and is reproducible from the repository.

Route	Recall@4 (canon. / adv.)	MRR (canon. / adv.)
Dense only	100.0% / 94.1%	0.978 / 0.843
Hybrid + RRF	100.0% / 100.0%	1.000 / 0.873
Hybrid + RRF + rerank	100.0% / 94.1%	1.000 / 0.941

falls back to local JSONL tracing, and web search falls back across four providers. Degradation is a design decision, not an incident.

3 Hybrid Retrieval and the Cost of Thresholds

The retrieval pipeline runs two parallel recall paths — a BM25 lexical path over jieba-tokenized chunks and a dense path over Chroma-stored embeddings — and merges them with reciprocal rank fusion using $k = 60$ [2]. RRF is chosen over score interpolation because BM25 scores are unbounded while cosine similarities lie in $[0, 1]$; rank-based fusion makes the merge invariant to these incompatible scales and requires no tuning.

Merged candidates are capped at ten and passed to a cross-encoder reranker (gte-rerank-v2 [7]). Because relevance scores from the reranker are calibrated probabilities, the pipeline applies a two-part filter: an absolute floor (0.05, or a 0.08 floor on the top hit, below which the system escalates to refusal) and a relative filter at 15% of the head score.

Table 1 reports the fixed-benchmark comparison. Three observations matter. First, canonical questions are essentially saturated by modern embeddings — differentiation only appears on the adversarial set, which contains colloquial, code-mixed, and abbreviation-heavy phrasings. Second, RRF raises adversarial Recall@4 from 94.1% to 100.0%: BM25 supplies exactly the literal matching that dense embeddings miss on jargon. Third, adding the reranker improves MRR from 0.873 to 0.941 — the correct chunk is ranked markedly higher — while recall drops by one borderline hit, which the relative threshold removes. We retain the threshold deliberately: downstream, the evidence-grounded refusal path prefers an honest “not in the knowledge base” over an answer woven from a weakly related chunk. The trade-off is explicit, benchmarked, and tunable per deployment.

Two engineering details extend the result to multi-knowledge-base deployments. Recall paths run per knowledge base without reranking; merged, deduplicated candidates are then reranked exactly once, cutting reranker API cost by a factor of the base count while keeping scores comparable across bases. And because the reranker is a remote API, failures fall back to RRF ordering, with a circuit breaker (three consecutive failures trigger a 60s cooldown) so that endpoint outages degrade ranking quality rather than availability.

4 A Determinism-Keyed Cache Hierarchy

Per-request latency is dominated by two remote calls: one embedding of the query and one rerank of the candidates. Both are deterministic functions of their inputs for a fixed model, which makes them safely cacheable at the query level — the cache changes *whether* a value is recomputed, never *what* the value is.

DocMind therefore maintains four layers, ordered by where determinism is easiest to argue: a CPU-level tokenizer cache over unchanged chunks; a persistent SQLite embedding cache over document chunks (keyed by text hash, invalidated on write failure); an in-process LRU over query embeddings; and an in-process LRU over rerank results. Cache keys embed the model name and — for the rerank cache — a SHA-256 fingerprint of the candidate set, so that a

Table 2: Same-code A/B with query-level caches toggled by flag: the identical 12-question set replayed at three concurrency levels. Warm passes skip both upstream round-trips; the answer-level semantic cache (similarity-matched, threshold 0.92) sits above this layer and is unaffected.

Scenario	Baseline (QPS / P95)	With caches (QPS / P95)
Concurrency 1 — first pass, cold	2.46 / 1,040 ms	2.84 / 1,059 ms
Concurrency 4 — second pass, warm	2.97 / 1,390 ms	741 / 6 ms
Concurrency 8 — third pass, warm	3.10 / 2,573 ms	801 / 11 ms

Table 3: Recall@4 of retrieval after rewriting, baseline (untuned 1.5B) versus LoRA-tuned rewriter, on 126 deterministic samples.

Group (n)	Baseline rewrite	Tuned rewrite	Delta
Overall (126)	0.754	0.794	+3.97 pp
Canonical (81)	0.827	0.889	+6.17 pp
Adversarial (45)	0.622	0.622	0.00

model switch or a knowledge-base update changes the key and old entries expire structurally, with no manual invalidation. Only successful results are cached, so failure and circuit-breaker semantics are unchanged; notably, hot queries served from cache remain available even while the breaker is open.

Table 2 reports a same-code A/B evaluation on the open retrieval endpoint. The same 12-question set is replayed at three concurrency levels, and because the caches persist across passes the first pass is entirely cold — every question traverses the full pipeline (12 misses) — while the second and third passes are served entirely from cache (24 hits). Replaying a warm question set models the repeat-query pattern that dominates FAQ-shaped enterprise traffic. Cache toggles are environment flags, so both runs execute identical code. Warm-pass throughput improves by roughly two orders of magnitude and P95 collapses from 2,573 ms to 11 ms; the cold pass is statistically unchanged, demonstrating that the caches neither help nor hurt long-tail queries. Hit rates are exported as Prometheus counters (67%: 12 misses, 24 hits).

5 Narrow Fine-Tuning: a LoRA Query Rewriter

Colloquial phrasings are the weak spot of dense recall. The common fix — asking the chat LLM to rewrite each query — costs one extra cloud round-trip per turn. DocMind instead fine-tunes a 1.5B model (Qwen2.5-1.5B-Instruct [6]) as a dedicated rewriter: a rank-8 LoRA adapter [4] trained with LLaMA-Factory on a MacBook (about 20 minutes), merged and served locally through Ollama as GGUF. The rewriter is a bypass component: if it fails or outputs nothing, retrieval proceeds with the raw query.

Training data is constructed deterministically: the 47 benchmark seed questions are expanded by fixed-seed noise rules into colloquial, code-mixed, typo, and preamble variants — eight per question plus the canonical form — yielding 364 deduplicated pairs with a 90/10 train-validation split. The A/B evaluation draws from a disjoint noise-seed range: three held-out variants per question (141 generated, 126 valid after deduplication), so no evaluation string appears in the training corpus. Because the seeds and rules are fixed, any rerun reproduces the identical datasets. On this evaluation, the tuned rewriter improves overall Recall@4 from 0.754 to 0.794 (+3.97 points) and canonical-group recall from 0.827 to 0.889 (+6.17 points), at a local rewrite latency of about 400 ms (Table 3). The rewrite step itself now costs zero API calls.

The adversarial group did not move, and a controlled experiment explains why: retrieving

with the original canonical adversarial questions — bypassing the rewriter entirely — yields only 0.588. The ceiling on this corpus is corpus coverage, not rewriting quality; indeed the tuned rewriter’s output retrieves better (0.622) than the human-written canonical phrasing itself (0.588). The correct next investment is therefore corpus expansion, not more training. We consider this attribution experiment the most transferable lesson of the study: when an A/B plateaus, test the ceiling before tuning the model.

6 Cost-Aware Model Routing

DocMind routes requests between cloud models and locally hosted ones through a pure-function decision layer with a fixed priority: explicit model selection, multimodal input, tool use, and chain-of-thought requests stay on the cloud model; greeting-class and short requests are routed to a local endpoint when configured; a deterministic FAQ-split experiment hashes the question into buckets so that identical questions always take the same route. When the local endpoint is the primary target, a cloud fallback is pre-registered, and in streaming scenarios the switch happens only before the first token is produced — a stream never changes backend mid-answer. Routing decisions are exported as Prometheus counters labeled by backend and reason, and the cost effect is recomputed from logs by a report script rather than asserted.

7 RetrievalOps and Governance

The platform treats retrieval quality as an operational property. A benchmark harness and an in-product evaluation laboratory expose per-stage retrieval timings and miss analysis; RAGAS-style generation metrics [5] (faithfulness, answer relevance, context precision and recall) are computed by an LLM judge with per-metric fault isolation. Negative feedback flows into a badcase queue and, from there, into the benchmark set — the same closed loop that produced the attribution study in Section 5.

Governance covers the enterprise baseline: hashed API keys with scope, rotation, and expiry; an immutable audit log with CSV export; alert rules with deduplication and a 24-hour cost cap; hot-backup of the SQLite store via `VACUUM INTO`; and per-key rate limiting on the open retrieval endpoint. The same endpoint is additionally exposed as an MCP tool server, so that coding assistants such as Claude Desktop or Cursor can query the knowledge base under the server-side access-control policy. During a 72-hour observation window the platform answered 1,126 calls with 94.85% availability at an average cost of about 4.4 CNY per thousand calls, with all failures isolated to auxiliary sub-tasks rather than the answering path.

8 Limitations

The evaluation corpus is deliberately small (47 seed questions, one 12-chunk example knowledge base), which bounds what any fine-tuning can show and makes absolute numbers illustrative rather than general. The rewriting A/B shares its canonical questions with training by construction — held-out noise seeds vary the surface forms, not the questions — so the measured gains are in-distribution rewriting improvements rather than evidence of generalization to unseen topics. The hot-workload cache gains are workload-shaped: deployments with uniformly unique queries will see cold-path behavior, not the 801 QPS tail. The platform targets single-node self-hosting on SQLite and Chroma; horizontal scaling, multi-tenancy, and connector-based corpus synchronization (e.g. wiki and drive sync with permission mirroring) are open work. Finally, the LLM-judge metrics inherit the usual judge biases and are reported alongside, not instead of, retrieval-level ground truth.

9 Conclusion

DocMind demonstrates that the distance between a RAG demo and a RAG system is measurable and bridgeable: a fixed benchmark turns retrieval design into engineering; determinism-keyed caches turn cost and latency into controlled quantities; narrow LoRA fine-tuning turns a per-request cloud dependency into a local asset with a quantified payoff; and a governance layer turns all of it into something an organization can actually run. The system is deliberately boring where it can be — owned code, graceful degradation, reproducible numbers — which is precisely what makes its few novel choices, such as the determinism-keyed rerank cache and the attribution-first fine-tuning workflow, easy to trust.

References

- [1] P. Lewis, P. Perez, E. Piktus, et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. *NeurIPS*, 2020.
- [2] G. V. Cormack, C. L. A. Clarke, and S. Büttcher. Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. *SIGIR*, 2009.
- [3] S. Robertson and H. Zaragoza. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*, 2009.
- [4] E. J. Hu, Y. Shen, P. Wallis, et al. LoRA: Low-rank adaptation of large language models. *ICLR*, 2022.
- [5] S. Es, J. James, L. Espinosa-Anke, and S. Schockaert. RAGAS: Automated evaluation of retrieval augmented generation. *EACL (System Demonstrations)*, 2024.
- [6] A. Yang, B. Yang, B. Zhang, et al. Qwen2.5 technical report. *arXiv:2412.15115*, 2024.
- [7] Z. Li, X. Zhang, Y. Zhang, et al. Towards general text embeddings with multi-stage contrastive learning. *arXiv:2308.03281*, 2023.
- [8] Y. Gao, Y. Xiong, G. Gao, et al. Retrieval-augmented generation for large language models: A survey. *arXiv:2312.10997*, 2023.